

Taints and Tolerations in Kubernetes

- [Overview](#)
- [About Taints and Tolerations](#)
- [Example](#)
- [Tolerations](#)
- [Related articles](#)

Overview

Platform	Azure
Owner of this SOP:	Fully Managed POD-B
Cloud Services	AKS

Tester	Reviewer
@ Sandeep Naidu Pinnenti	

About Taints and Tolerations

Taints allow a node to repel a set of pods. Tolerations allow the scheduler to schedule pods with matching taints. Taints and tolerations help prevent your pods from scheduling to undesirable nodes.

Taints and tolerations work together to ensure that pods are not scheduled onto inappropriate nodes. One or more taints are applied to a node; this marks that the node should not accept any pods that do not tolerate the taints.

In some cases, we do not want application pods to be placed on nodes that runs system or administration related task. In such scenarios we can taint a node so that pod that do not tolerate the taint cannot be placed in that node.

Example

List the nodes that are present in the cluster.

```
C:\Users\podbadmin>kubectl get nodes
NAME                                STATUS    ROLES    AGE     VERSION
aks-default-25191975-vmss00000d    Ready    agent    6h19m   v1.24.6
aks-podbttesting-20505491-vmss000008  Ready    agent    6h19m   v1.24.6
aks-podbttesting-20505491-vmss000009  Ready    agent    6h19m   v1.24.6

C:\Users\podbadmin>
```

Check the taints that are present on the node. One way is to use **kubectl describe node** command.

```
C:\Users\podbadmin>kubectl describe node aks-default-25191975-vmss00000d
Name:          aks-default-25191975-vmss00000d
Roles:         agent
Labels:        agentpool=default
               beta.kubernetes.io/arch=amd64
               beta.kubernetes.io/instance-type=Standard_D4s_v3
               beta.kubernetes.io/os=linux
               failure-domain.beta.kubernetes.io/region=australiaeast
```

Taints:
Unschedulable

Add a taint to a node using `kubectl taint`.

```
kubectl taint nodes node1 key1=value1:Effect
```

Below, we have placed a taint on node **aks-default-25191975-vmss00000d** with effect NoSchedule.

```
kubectl taint nodes node1 key1=value1:NoSchedule
```

The taint has key **nodepool**, value **system**, and taint effect **NoSchedule**. This means that no pod will be able to schedule onto “**aks-default-25191975-vmss00000d**” unless it has a matching toleration.

This means that pod will be evicted from the node (if it is already running on the node) and will not be scheduled onto the node (if it is not yet running on the node).

```
C:\Users\podbadmin>kubectl get nodes
NAME                                STATUS    ROLES    AGE      VERSION
aks-default-25191975-vmss00000d    Ready    agent    6h22m    v1.24.6
aks-podbtesting-20505491-vmss000008    Ready    agent    6h22m    v1.24.6
aks-podbtesting-20505491-vmss000009    Ready    agent    6h22m    v1.24.6

C:\Users\podbadmin>kubectl taint node aks-default-25191975-vmss00000d nodepool=system:NoSchedule
node/aks-default-25191975-vmss00000d tainted

C:\Users\podbadmin>
```

OK Administrator: Command Prompt

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
nginx-deploy-5c5d77d48d-7qvcn	1/1	Running	0	81s	10.244.1.6	aks-poddbtesting-20505491-vmss000008	<none>	<none>
nginx-deploy-5c5d77d48d-9ckjj	1/1	Running	0	81s	10.244.1.8	aks-poddbtesting-20505491-vmss000008	<none>	<none>
nginx-deploy-5c5d77d48d-g7drs	1/1	Running	0	81s	10.244.1.9	aks-poddbtesting-20505491-vmss000008	<none>	<none>
nginx-deploy-5c5d77d48d-jpcgc	1/1	Running	0	81s	10.244.2.34	aks-poddbtesting-20505491-vmss000009	<none>	<none>
nginx-deploy-5c5d77d48d-q2692	1/1	Running	0	81s	10.244.1.7	aks-poddbtesting-20505491-vmss000008	<none>	<none>
nginx-deploy-5c5d77d48d-xzax2	1/1	Running	0	81s	10.244.2.33	aks-poddbtesting-20505491-vmss000009	<none>	<none>

The above example used effect of `NoSchedule`. Alternatively, you can use effect of **`NoExecute`**. This means that pod will be evicted from the node (if it is already running on the node) and will not be scheduled onto the node (if it is not yet running on the node).

```
kubectl taint nodes node1 key1=value1:NoExecute
```

```
Administrator: Command Prompt
C:\Users\podbadmin>kubectl get nodes
NAME                                STATUS    ROLES    AGE      VERSION
aks-default-25191975-vmss00000d    Ready    agent    6h35m    v1.24.6
aks-podtesting-20505491-vmss000008  Ready    agent    6h35m    v1.24.6
aks-podtesting-20505491-vmss000009  Ready    agent    6h35m    v1.24.6

C:\Users\podbadmin>kubectl describe nodes aks-podtesting-20505491-vmss000008
Name: aks-podtesting-20505491-vmss000008
Roles: agent
Labels: agentpool=podtesting
       beta.kubernetes.io/arch=amd64
       beta.kubernetes.io/instance-type=Standard_D4s_v3
       beta.kubernetes.io/os=linux
       failure-domain.beta.kubernetes.io/region=australiaeast
       failure-domain.beta.kubernetes.io/zone=australiaeast-1
       kubernetes.azure.com/agentpool=podtesting
       kubernetes.azure.com/cluster=it-aeu1-npe-arg-podb-aks-nodes
       kubernetes.azure.com/kubelet-identity-client-id=09999fbef-96f7-45a8-8555-42e85192724e
       kubernetes.azure.com/node=user
       kubernetes.azure.com/node-image-version=AKSUbuntu-1804gen2containerd-2022.11.02
       kubernetes.azure.com/os-sku=Ubuntu
       kubernetes.azure.com/role=agent
       kubernetes.azure.com/storageprofile=managed
       kubernetes.azure.com/storage-tier=Premium_LRS
       kubernetes.io/arch=amd64
       kubernetes.io/hostname=aks-podtesting-20505491-vmss000008
       kubernetes.io/os=linux
       kubernetes.io/role=agent
       node-role.kubernetes.io/agent=
       node.kubernetes.io/instance-type=Standard_D4s_v3
       storageprofile=managed
       storage-tier=Premium_LRS
       topology.disk.csi.azure.com/zone=australiaeast-1
       topology.kubernetes.io/region=australiaeast
       topology.kubernetes.io/zone=australiaeast-1
Annotations: csi.volume.kubernetes.io/nodeid:
              {"disk.csi.azure.com":"aks-podtesting-20505491-vmss000008","file.csi.azure.com":"aks-podtesting-20505491-vmss000008"}
              node.alpha.kubernetes.io/ttl: 0
              volumes.kubernetes.io/controller-managed-attach-detach: true
CreationTimestamp: Thu, 01 Dec 2022 03:39:55 +0000
Taints: <none>

C:\Users\podbadmin>kubectl get nodes
NAME                                STATUS    ROLES    AGE      VERSION
aks-default-25191975-vmss00000d    Ready    agent    6h39m    v1.24.6
aks-podtesting-20505491-vmss000008  Ready    agent    6h39m    v1.24.6
aks-podtesting-20505491-vmss000009  Ready    agent    6h39m    v1.24.6

C:\Users\podbadmin>kubectl taint node aks-podtesting-20505491-vmss000008 nodepool=user:NoExecute
node/aks-podtesting-20505491-vmss000008 tainted

C:\Users\podbadmin>
```

Once the taint is applied to the node, all the pods that are present in the node are evicted and are scheduled in other node.

```
C:\Users\podbadmin>kubectl get pods -o wide
NAME                                READY    STATUS    RESTARTS  AGE   IP              NODE                                NOMINATED NODE    READINESS GATES
nginx-deploy-5c5d77d48d-5bp98      1/1      Running   0          15s   10.244.2.37     aks-podtesting-20505491-vmss000009  <none>             <none>
nginx-deploy-5c5d77d48d-872vp      1/1      Running   0          15s   10.244.2.35     aks-podtesting-20505491-vmss000009  <none>             <none>
nginx-deploy-5c5d77d48d-b6gdb      1/1      Running   0          15s   10.244.2.36     aks-podtesting-20505491-vmss000009  <none>             <none>
nginx-deploy-5c5d77d48d-cxfv1      1/1      Running   0          15s   10.244.2.38     aks-podtesting-20505491-vmss000009  <none>             <none>
nginx-deploy-5c5d77d48d-jpcgc      1/1      Running   0          15m   10.244.2.34     aks-podtesting-20505491-vmss000009  <none>             <none>
nginx-deploy-5c5d77d48d-zxqz2      1/1      Running   0          15m   10.244.2.33     aks-podtesting-20505491-vmss000009  <none>             <none>

C:\Users\podbadmin>
```

Tolerations

Tolerations allow the scheduler to schedule pods with matching taints.

```

C:\Users\podadmin>kubectl describe node aks-podtesting-20505491-vmss000009
Name: aks-podtesting-20505491-vmss000009
Roles: agent
Labels: agentpool=podtesting
        beta.kubernetes.io/arch=amd64
        beta.kubernetes.io/instance-type=Standard_D4s_v3
        beta.kubernetes.io/os=linux
        failure-domain.beta.kubernetes.io/region=australiaeast
        failure-domain.beta.kubernetes.io/zone=australiaeast-2
        kubernetes.azure.com/agentpool=podtesting
        kubernetes.azure.com/cluster=it-ae1-rpe-arg-pod-aks-nodes
        kubernetes.azure.com/kubelet-identity-client-id=0999fbef-96f7-45a8-8555-42e85192724e
        kubernetes.azure.com/node=user
        kubernetes.azure.com/node-image-version=AKSUbuntu-1804gen2containerd-2022.11.02
        kubernetes.azure.com/os=skuUbuntu
        kubernetes.azure.com/role=agent
        kubernetes.azure.com/storageprofile=managed
        kubernetes.azure.com/storagetier=Premium_LRS
        kubernetes.io/arch=amd64
        kubernetes.io/hostname=aks-podtesting-20505491-vmss000009
        kubernetes.io/os=linux
        kubernetes.io/role=agent
        node-role.kubernetes.io/agent=
        node.kubernetes.io/instance-type=Standard_D4s_v3
        storageprofile=managed
        storage-tier=Premium_LRS
        topology.disk.csi.azure.com/zone=australiaeast-2
        topology.kubernetes.io/region=australiaeast
        topology.kubernetes.io/zone=australiaeast-2
Annotations: csi.volume.kubernetes.io/nodeid:
              ("disk.csi.azure.com":"aks-podtesting-20505491-vmss000009","file.csi.azure.com":"aks-podtesting-20505491-vmss000009")
              node.alpha.kubernetes.io/ttl: 0
              volumes.kubernetes.io/controller-managed-attach-detach: true
CreationTimestamp: Thu, 01 Dec 2022 03:40:05 +0000
Taints: nodepool=user1:NoExecute
        nodepool=user1:NoSchedule
Unschedulable: false
Lease:
  HolderIdentity: aks-podtesting-20505491-vmss000009
  AcquireTime: <unset>
  RenewTime: Thu, 01 Dec 2022 10:54:55 +0000

```

Pods are in pending state as the toleration is not defined in POD YAML spec.

nginx-deploy-844f6c8c7f-1hxxvw	0/1	Pending	0	0s	<none>	<none>	<none>	<none>	<none>
nginx-deploy-844f6c8c7f-1hxxvw	0/1	Pending	0	0s	<none>	aks-podtesting-20505491-vmss000009	<none>	<none>	<none>
nginx-deploy-844f6c8c7f-26thz	0/1	Pending	0	0s	<none>	<none>	<none>	<none>	<none>
nginx-deploy-844f6c8c7f-hzjjk	0/1	Pending	0	0s	<none>	<none>	<none>	<none>	<none>
nginx-deploy-844f6c8c7f-26thz	0/1	Pending	0	0s	<none>	aks-podtesting-20505491-vmss000009	<none>	<none>	<none>
nginx-deploy-844f6c8c7f-hzjjk	0/1	Pending	0	0s	<none>	aks-podtesting-20505491-vmss000009	<none>	<none>	<none>

Add tolerations to the Pod YAML file that will tolerate the taint.

- key: "nodepool"
 - operator: "Equal"
 - value: "user1"
 - effect: "NoSchedule"
- key: "nodepool"
 - operator: "Equal"
 - value: "user1"
 - effect: "NoExecute"

Once the tolerations are added to the Pod YAML spec, we can see the pods are in running state.

nginx-deploy-7fd99c4ff8-kpf9k	0/1	Terminating	0	8m20s	<none>	<none>	<none>	<none>	<none>
nginx-deploy-844f6c8c7f-686h9	0/1	Pending	0	0s	<none>	<none>	<none>	<none>	<none>
nginx-deploy-844f6c8c7f-686h9	0/1	Pending	0	0s	<none>	aks-podtesting-20505491-vmss000009	<none>	<none>	<none>
nginx-deploy-844f6c8c7f-686h9	0/1	ContainerCreating	0	0s	<none>	aks-podtesting-20505491-vmss000009	<none>	<none>	<none>
nginx-deploy-844f6c8c7f-26thz	1/1	Running	0	5s	10.244.2.41	aks-podtesting-20505491-vmss000009	<none>	<none>	<none>
nginx-deploy-7fd99c4ff8-td95t	0/1	Terminating	0	8m21s	<none>	<none>	<none>	<none>	<none>
nginx-deploy-7fd99c4ff8-td95t	0/1	Terminating	0	8m21s	<none>	<none>	<none>	<none>	<none>
nginx-deploy-844f6c8c7f-pkhpz	0/1	Pending	0	0s	<none>	<none>	<none>	<none>	<none>
nginx-deploy-844f6c8c7f-pkhpz	0/1	Pending	0	0s	<none>	aks-podtesting-20505491-vmss000009	<none>	<none>	<none>
nginx-deploy-844f6c8c7f-pkhpz	0/1	ContainerCreating	0	0s	<none>	aks-podtesting-20505491-vmss000009	<none>	<none>	<none>
nginx-deploy-844f6c8c7f-1hxxvw	1/1	Running	0	7s	10.244.2.39	aks-podtesting-20505491-vmss000009	<none>	<none>	<none>
nginx-deploy-84c446c78f-2b6w2	0/1	Terminating	0	7m20s	<none>	<none>	<none>	<none>	<none>
nginx-deploy-84c446c78f-2b6w2	0/1	Terminating	0	7m20s	<none>	<none>	<none>	<none>	<none>
nginx-deploy-844f6c8c7f-bzk5q	0/1	Pending	0	0s	<none>	<none>	<none>	<none>	<none>
nginx-deploy-844f6c8c7f-bzk5q	0/1	Pending	0	0s	<none>	aks-podtesting-20505491-vmss000009	<none>	<none>	<none>
nginx-deploy-844f6c8c7f-bzk5q	0/1	ContainerCreating	0	0s	<none>	aks-podtesting-20505491-vmss000009	<none>	<none>	<none>
nginx-deploy-844f6c8c7f-686h9	1/1	Running	0	5s	10.244.2.42	aks-podtesting-20505491-vmss000009	<none>	<none>	<none>
nginx-deploy-84c446c78f-hlpsq	0/1	Terminating	0	7m22s	<none>	<none>	<none>	<none>	<none>
nginx-deploy-84c446c78f-hlpsq	0/1	Terminating	0	7m22s	<none>	<none>	<none>	<none>	<none>
nginx-deploy-844f6c8c7f-pkhpz	1/1	Running	0	6s	10.244.2.43	aks-podtesting-20505491-vmss000009	<none>	<none>	<none>
nginx-deploy-84c446c78f-zgntn	0/1	Terminating	0	7m24s	<none>	<none>	<none>	<none>	<none>
nginx-deploy-84c446c78f-zgntn	0/1	Terminating	0	7m24s	<none>	<none>	<none>	<none>	<none>
nginx-deploy-844f6c8c7f-bzk5q	1/1	Running	0	5s	10.244.2.44	aks-podtesting-20505491-vmss000009	<none>	<none>	<none>

```
C:\Users\podbadmin>kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
nginx-deploy-844f6c8c7f-26thz	1/1	Running	0	24m	10.244.2.41	aks-podtesting-20505491-vmss000009	<none>	<none>
nginx-deploy-844f6c8c7f-686h9	1/1	Running	0	24m	10.244.2.42	aks-podtesting-20505491-vmss000009	<none>	<none>
nginx-deploy-844f6c8c7f-bzk5q	1/1	Running	0	24m	10.244.2.44	aks-podtesting-20505491-vmss000009	<none>	<none>
nginx-deploy-844f6c8c7f-hzjjk	1/1	Running	0	24m	10.244.2.40	aks-podtesting-20505491-vmss000009	<none>	<none>
nginx-deploy-844f6c8c7f-lhxvw	1/1	Running	0	24m	10.244.2.39	aks-podtesting-20505491-vmss000009	<none>	<none>
nginx-deploy-844f6c8c7f-pkhpz	1/1	Running	0	24m	10.244.2.43	aks-podtesting-20505491-vmss000009	<none>	<none>

Related articles

- <https://kubernetes.io/docs/concepts/scheduling-eviction/taint-and-toleration/>